



Education

2020-2024 Chongqing Normal University (undergraduate student)

Intern

2024-2025 The Hong Kong University of Science and Technology (Guangzhou) (research assistant)

2026-2026 LLM Startup (research intern for post-training and on-policy distillation)

2026-2026 Hunyuan Qingyun Project (foundational reinforcement learning team)

Current Interests

◦ Reinforcement Learning with Verifiable Rewards ◦ On-Policy Distillation

Research

2026 **NeurIPS submission (4, 4, 4).** Zhengpeng Xie, Li Lyna Zhang, Zeke Xie, Mao Yang *et al.* *Trust Region Policy Distillation.* [Link](#)

This work presents TOP-D, which reformulates on-policy distillation by distilling from a dynamically constructed proximal teacher rather than the original teacher, enabling more stable and efficient training.

2025 **ICML 2025 accept.** Zhengpeng Xie*, Qiang Zhang*, Fan Yang*, Marco Hutter, Renjing Xu *et al.* *Simple Policy Optimization.* [Link](#)

This work proposes a new model-free policy gradient algorithm called Simple Policy Optimization (SPO), a powerful algorithm adapted by the state-of-the-art VLA model $\pi_{0.6}^$ and FPO++.*

2026 **CVPR 2026 accept.** Zening Sun*, Zhengpeng Xie*, Lichen Bai, Shitong Shao, Shuo Yang, Zeke Xie *et al.* *CRAFT: Aligning Diffusion Models with Fine-Tuning Is Easier Than You Think.* [Link](#)

This work proposes a lightweight diffusion model fine-tuning framework, CRAFT, which can be theoretically understood as optimizing a lower bound of group-based reinforcement learning.

2026 **arXiv preprint.** Zhengpeng Xie, Jiahang Cao, Qiang Zhang, Jianxiong Zhang, Changwei Wang, Renjing Xu *et al.* *Representation Convergence: Mutual Distillation is Secretly a Form of Regularization.* [Link](#)

This work provides a formal proof for a long-standing assumption that policy robustness facilitates generalization, and further empirically demonstrates that online mutual distillation helps improve generalization performance.

2025 **arXiv preprint.** Zhengpeng Xie*, Yulong Zhang* *et al.* *A Dual-Agent Adversarial Framework for Robust Generalization in Deep Reinforcement Learning.* [Link](#)

This work proposes an interesting adversarial framework for improving the generalization performance of reinforcement learning policies.

2025 **ICML 2025 workshop accept.** Junbin Qiu, Zhengpeng Xie, Xiangda Yan, Yongjie Yang, Yao Shu *et al.* *Zeroth-Order Optimization is Secretly Single-Step Policy Optimization.* [Link](#)

This work bridges the gap between policy optimization and zeroth-order optimization techniques, transfers the variance reduction technique from reinforcement learning to zeroth-order optimization.

Link

◦ Google Scholar: [Click here](#)

◦ GitHub: [Click here](#)